

Project 2
Database Management Systems (DM556)



**UNIVERSITY OF
SOUTHERN DENMARK**

Group 2
Mark Jervelund (Mjerv15) Troels B. Petersen
(trpet15) IMADA

April 4, 2017

Overall Status

The group managed to complete the tasks and therefore the project is considered complete.

Division of Labor

The group worked on the project either sitting together at the university or at home remotely working together and splitting tasks when possible. A lot of the time was spent understanding how to implement a solution. Especially Sort and Merge-Join was not very straight forward. The work was very evenly divided - both when writing the code, but also when writing the report.

Specification

The group was tasked with implementing four operators; Selection, Projection, Sort and Merge-Join.

Selection

Selection is a very basic operation in database management. It uses relational algebra to select the elements. Now the spec for this operator says that every query is combined with a relational *or*. This means that nothing "fancy" has to be done. It should simply select everytime one or more queries return true.

Projection

The projection is also one of the more basic operations in database management. A projection extracts the columns from a relation, however, unlike in relational algebra, this operator does not eliminate duplicates.

Sort

Sort has to be external. External sorting is used in applications where huge amounts of data has to be sorted, thus the data has to be sorted in chunks since it cannot all be in main memory. A variant of merge-sort will be used, since it can sort on parts of the data and then combine the sorted parts.

Merge-Join

The merge-join assumes that the both inputs are sorted. It then has to merge where possible.

Implementation

Selection.java

Selection starts by assigning local protected variables some values from the parameters.

```
16  */
    public Selection(Iterator iter, Predicate... preds) {
        this.iterator = iter;
        this.predicates = preds;
20  this.schema = iter.schema;
        this.tuple = null;
```

The selection process takes place in the hasNext() function. Here it keeps checking if there are more elements to be selected using the evaluate() function. It returns true if it allows the selection and false if there are no more elements to be selected.

```
66  * Returns true if there are more tuples, false otherwise.
    */
    public boolean hasNext() {
        while(this.iterator.hasNext()) {
```

```

70      this.tuple = this.iterator.getNext();

      for(int i = 0; i < this.predicates.length; ++i) {
          if(this.predicates[i].evaluate(this.tuple)) {
75              return true;
          }
      }
  }

```

The getNext() function is what actually gets the elements. It returns a tuple containing the next element. If there are no more elements to be selected, it will throw an exception, telling that there are no more tuples.

```

83      * Gets the next tuple in the iteration.
      *
85      * @throws IllegalStateException if no more tuples
      */
      public Tuple getNext() {
          if(this.tuple == null) {
              throw new IllegalStateException("no_more_tuples");
90          } else {
              Tuple tuple = this.tuple;

```

Projection.java

Sort.java

MergeJoin.java

The merge join function first merges the two schemas

```

41      public MergeJoin(Iterator left , Iterator right , Integer lcol , Integer
      ↪ rcol) {
          this.left = left;
          this.right = right;
          this.lcol = lcol;
45          this.rcol = rcol;
          schema = Schema.join(left.schema , right.schema);

```

The Hasnext function then first selects a left tuple, and compares it to all right tuples, if it two that are comparable it returns true. else it loops over all combinations and returns false if none are found. but also the next element in the array is stored in the next variable

```

88      public boolean hasNext() {
          while (true) {
90              if (outer == null) {
                  if (left.hasNext()) {
                      outer = left.getNext();
                  } else { return false; }
              }
95              while (this.right.hasNext()) {
                  Tuple rightTuple = right.getNext();
                  next = Tuple.join(outer , rightTuple , schema);
                  if (outer.getField(lcol) == rightTuple.getField(rcol)) {
                      return true;
100              }
                  }
              outer = null;
              right.restart();

```

The next variable is then stored in a temp variable, set to null, and returned from the temp variable

```

111     public Tuple getNext() {
        // validate the next tuple
        if (next == null) {
            throw new IllegalStateException("no_more_tuples");
115     }
        // return (and forget) the tuple
        Tuple tuple = next;
        next = null;
        return tuple;
120     }

```

Testing

Testing this time around was very successful. It reports that test1, test2 and test3 completed successfully.

Further more when comparing the expected output with the supplied ExpectedOutput.txt we got the same output except for some initializing prints and some lines from the explain print statement.

Appendix

Selection.java

```

1  package relop;

    /**
     * The selection operator specifies which tuples to retain under a
     *   ↪ condition; in
5   * Minibase, this condition is simply a set of independent predicates
     *   ↪ logically
     * connected by OR operators.
     */
    public class Selection extends Iterator {

10     protected Iterator iterator;
        protected Predicate[] predicates;
        protected Tuple tuple;
        /**
15     * Constructs a selection, given the underlying iterator and predicates.
     *
     */
        public Selection(Iterator iter, Predicate... preds) {
            this.iterator = iter;
            this.predicates = preds;
20     this.schema = iter.schema;
            this.tuple = null;
        }

        /**
25     * Gives a one-line explanation of the iterator, repeats the call on any
     * child iterators, and increases the indent depth along the way.
     */
        public void explain(int depth) {
30     //this.schema(depth);
            System.out.print("Selection_:");

            for(int i = 0; i < this.predicates.length - 1; ++i) {

```

```

        System.out.print(this.predicates[i].toString() + "_OR_");
    }
35
    System.out.println(this.predicates[this.predicates.length - 1]);
    this.iterator.explain(depth + 1);
}

40
/**
 * Restarts the iterator, i.e. as if it were just constructed.
 */
public void restart() {
    this.iterator.restart();
45
    this.tuple = null;
}

/**
 * Returns true if the iterator is open; false otherwise.
50
 */
public boolean isOpen() {
    return this.iterator != null;
}

55
/**
 * Closes the iterator, releasing any resources (i.e. pinned pages).
 */
public void close() {
    if(this.iterator != null) {
60
        this.iterator.close();
        this.iterator = null;
    }
}

65
/**
 * Returns true if there are more tuples, false otherwise.
 */
public boolean hasNext() {
    while(this.iterator.hasNext()) {
70
        this.tuple = this.iterator.getNext();

        for(int i = 0; i < this.predicates.length; ++i) {
            if(this.predicates[i].evaluate(this.tuple)) {
75
                return true;
            }
        }

        return false;
80
    }

    /**
     * Gets the next tuple in the iteration.
     *
     * @throws IllegalStateException if no more tuples
85
     */
    public Tuple getNext() {
        if(this.tuple == null) {
            throw new IllegalStateException("no_more_tuples");
90
        } else {

```

```

        Tuple tuple = this.tuple;
        this.tuple = null;
        return tuple;
    }
}
} // public class Selection extends Iterator

```

sort.java

```

1  package relop;

import global.*;
import heap.HeapFile;

5  import java.io.File;
import java.util.ArrayList;
import java.util.Arrays;
import java.util.HashMap;

10 public class Sort extends Iterator implements GlobalConst {
    protected Iterator iterator;
    protected HeapFile file;
    protected FileScan scan;

15
    /**
     * Constructs a sort operator.
     * @param sortMemSize the size the memory used for internal sorting. For
     *   ↪ simplicity, you can assume it is in the unit of tuples.
     * @param bufSize the total buffer size for the merging phase in the unit
     *   ↪ of page.
20
     *
     */
    public Sort(Iterator iter, int sortfield, int sortMemSize, int bufSize) {

25
        this.iterator = iter;
        schema = iter.schema;
        HeapFile[] records = new HeapFile[bufSize];
        Tuple[] internal = new Tuple[sortMemSize];
        Object[] all = new Object[sortMemSize];
        HashMap<Object, Tuple> hashmap = new HashMap<Object, Tuple>();

30
        int pos = 0;
        // read data into sorting area
        while (iter.hasNext()) {
            // Load the records into the internal memory
35
            for (int i = 0; i < sortMemSize; i++) {
                if (iter.hasNext()) {
                    internal[i] = iter.getNext();
                    all[i] = internal[i].getField(0);
                    hashmap.put(all[i], internal[i]);
40
                }
            } // for

            ArrayList<Object> queue = new ArrayList<Object>();

45
            for (Object object : all) {
                if (object != null) {
                    queue.add(object);
                }
            }
        }
    }
}

```

```

    }
    }

50     all = queue.toArray();
    // Sort the tuples
    java.util.Arrays.sort(all);

55     records[pos] = new HeapFile(null);
    for (Object object : all) {
        records[pos].insertRecord(hashmap.get(object).
            ↪ data);
    }
    pos++;
60 } // PASS 1

    file = sorter(records, bufSize, iter, sortfield)[0];
    scan = new FileScan(iter.schema, file);
}

65 private HeapFile[] sorter(HeapFile[] records, int bufSize, Iterator
    ↪ iter, int sortfield) {

    int heapCount = getHeaps(records);
    if (heapCount == 1) {
70         return records;
    }
    if (heapCount >= bufSize) {
        heapCount = bufSize - 1;
    }
75 FileScan[] scan = new FileScan[heapCount];

    // Create a new filescan on every record in the current
    ↪ record array
    for (int i = 0; i < heapCount; i++) {
        scan[i] = new FileScan(iter.schema, records[i]);
80     }

    HeapFile file = new HeapFile(null);
    Tuple[] tuples = new Tuple[heapCount];
    int compared = 0;

85     // Load the tuples from the filescanner
    for (int i = 0; i < tuples.length; i++) {
        tuples[i] = scan[i].getNext();
    }

90     while (compared != heapCount) {
        Object[] smallest = { null, null };
        int smallestPos = 0;
        int current = 0;
95         for (Tuple tuple : tuples) {
            Object next = tuple.getField(sortfield);
            if (smallest[0] == null) {
                smallest[0] = next;
                smallest[1] = next;
            } else { // compare
100                 smallest[1] = next;
            }
        }
    }
}

```

```

105         java.util.Arrays.sort(smallest);
        if (smallest[0] == next) {
            smallestPos = current;
        }
        current++;
110     }

    file.insertRecord(tuples[smallestPos].data);
    if (scan[smallestPos].hasNext()) {
        tuples[smallestPos] = scan[smallestPos].
            ↪ getNext();
115     } else {
        tuples[smallestPos].setField(sortfield,
            ↪ Integer.MAX_VALUE);
        compared++;
    }

120     }
    records[heapCount - 1] = file;
    HeapFile[] rest = Arrays.copyOfRange(records, heapCount -
        ↪ 1, records.length);
    return sorter(rest, bufSize, iter, sortfield);
}

125 private int getHeaps(HeapFile[] records) {
    int result = 0;
    while (records[result] != null) {
        result++;
130     }
    return result;
}

@Override
135 public void explain(int depth) {
    FileScan fs = new FileScan(iterator.schema, file);
    fs.explain(depth);
}

@Override
140 public void restart() {
    scan.restart();
}

@Override
145 public boolean isOpen() {
    return scan.isOpen();
}

@Override
150 public void close() {
    if (scan != null){
        scan.close();
        scan = null;
155     }
}

```



```

    @Override
    public boolean hasNext() {
160         return scan.hasNext();
    }

    @Override
    public Tuple getNext() {
165         return scan.getNext();
        // throw new UnsupportedOperationException("Not implemented");
    }
}

```

MergeJoin.java

```

1 package relop;

import java.util.IllegalFormatException;
5
public class MergeJoin extends Iterator {

    /**
10     * The underlying left iterator.
    */
    protected Iterator left;

    /**
15     * The underlying right iterator.
    */
    protected Iterator right;

    /**
20     * left col.
    */
    protected Integer lcol;

    /**
25     * right col.
    */
    protected Integer rcol;

    /**
30     * Current tuple from left iterator.
    */
    protected Tuple outer;

    /**
35     * Next tuple to return.
    */
    protected Tuple next;

40
    public MergeJoin(Iterator left, Iterator right, Integer lcol, Integer
        ↪ rcol) {
        this.left = left;
        this.right = right;
    }
}

```

```

45      this.lcol = lcol;
      this.rcol = rcol;
      schema = Schema.join(left.schema, right.schema);

    }

50

    @Override
    public void explain(int depth) {

        indent(depth);
55      System.out.print("Merge_join_: ");
      for (int i = 0; i < this.schema.names.length - 1; i++) {
          System.out.println("{ " + this.schema.names[i] + " }");
      }
      System.out.println("{ " + this.schema.names[this.schema.names.length
        ↪ - 1] + " }");
60      this.left.explain(depth + 1);
      this.right.explain(depth + 1);
    }

    @Override
65    public void restart() {
        left.restart();
        right.restart();
        outer = null;
        next = null;

70    }

    @Override
    public boolean isOpen() {
        return (left != null);
75    }

    @Override
    public void close() {
        if (left != null) {
80            left.close();
            right.close();
            left = null;
            right = null;
        }

85    }

    @Override
    public boolean hasNext() {
        while (true) {
90            if (outer == null) {
                if (left.hasNext()) {
                    outer = left.getNext();
                } else { return false; }
            }
95            while (this.right.hasNext()) {
                Tuple rightTuple = right.getNext();
                next = Tuple.join(outer, rightTuple, schema);
                if (outer.getField(lcol) == rightTuple.getField(rcol)) {
                    return true;
100            }
        }
    }

```

```

    }
    outer = null;
    right.restart();
}
105 //          throw new IllegalStateException("debugging crash");

}

110 @Override
    public Tuple getNext() {
        // validate the next tuple
        if (next == null) {
            throw new IllegalStateException("no_more_tuples");
115         }
        // return (and forget) the tuple
        Tuple tuple = next;
        next = null;
        return tuple;
120     }
}

```

Test program output

```

1  Creating database...
   Replacer: Clock

   Running basic relational operator tests...
5  Test 1: Primitive relational operators

   ~> test selection (Age > 65 OR Age < 15)...

10 Selection : {3} > 65.0 OR {3} < 15.0
    FileScan : null
    DriverId  FirstName          LastName          Age          NumSeats
    -----
15  1          f1                  l1               7.7          101
    9          f9                  l9               69.3         109
    10         f10                 l10              77.0         110

   ~> test projection (columns 3 and 1)...

20 Projection : {3}, {1}
    FileScan : null
    Age       FirstName
    -----
25  7.7        f1
    15.4       f2
    23.1       f3
    30.8       f4
    38.5       f5
    46.2       f6
30  53.9       f7
    61.6       f8
    69.3       f9
    77.0       f10

```

35 ~> selection and projection (pipelined)...

Projection : {3}, {1}
 Selection : {3} > 65.0 OR {3} < 15.0
 FileScan : **null**

40 Age FirstName

7.7	f1
69.3	f9
77.0	f10

45

Test 1 completed without exception.

50

	Reads	Writes	Allocs	Pinned
insert	0	8	7	0
select	0	0	0	0
project	0	0	0	0
both	0	0	0	0

55

Test 2: Sorting Test

60 ... Inserted

~> sort numbers

checking the result.

65

Test 2 completed without error.

Test 3: MergeJoin operator

70

Projection : {DriverId}
 {FirstName}
 {LastName}
 {Age}
 {NumSeats}
 {DriverId}
 {GroupId}
 {FromDate}
 {ToDate}

75

FileScan : **null**

FileScan : **null**

80

DriverId	FirstName	LastName	Age	NumSeats
↪ DriverId	GroupId	FromDate	ToDate	

1	↪	Ahmed	Elmagarmid	25.0	5	1
85	1	↪	2	2/12/2006	2/14/2006	
1	↪	Ahmed	Elmagarmid	25.0	5	1
2	↪	3	2/15/2006	2/16/2006		
2	↪	Walid	Aref	27.0	13	2
2	↪	6	2/17/2006	2/20/2006		
2	↪	Walid	Aref	27.0	13	2

90	3	↪	7	2/18/2006	2/23/2006				
			Christopher	Clifton	18.0	4	3		
		↪	5	2/10/2006	2/13/2006				
	3		Christopher	Clifton	18.0	4	3		
		↪	4	2/18/2006	2/19/2006				
	3		Christopher	Clifton	18.0	4	3		
		↪	2	2/24/2006	2/26/2006				
	4		Sunil	Prabhakar	22.0	7	4		
		↪	1	2/19/2006	2/19/2006				
	5		Elisa	Bertino	26.0	5	5		
95		↪	7	2/14/2006	2/18/2006				
	6		Susanne	Hambruch	23.0	3	6		
		↪	6	2/25/2006	2/26/2006				
	8		Arif	Ghafoor	20.0	5	8		
		↪	5	2/20/2006	2/22/2006				
	9		Jeff	Vitter	19.0	10	9		
		↪	1	2/15/2006	2/15/2006				
	Test 3 completed without exception.								
	100	<hr/>							
			Reads	Writes	Allocs	Pinned			
<hr/>									
driver2		0	3	2	0				
105	rides2	0	3	2	0				
	m_join	0	0	0	0				
	<hr/>								
	All basic relational operator tests completed; verify output for ↪ correctness.								
110	Process finished with exit code 0								