

Project 2
Database Management Systems (DM556)



**UNIVERSITY OF
SOUTHERN DENMARK**

Group 2
Mark Jervelund (Mjerv15) Troels B. Petersen
(trpet15) IMADA

April 3, 2017

Overall Status

The group managed to complete the tasks and therefore the project is considered complete.

Division of Labor

The group worked on the project either sitting together at the university or at home remotely working together and splitting tasks when possible. A lot of the time was spent understanding how to implement a solution. Especially Sort and Merge-Join was not very straight forward. The work was very evenly divided - both when writing the code, but also when writing the report.

Specification

The group was tasked with implementing four operators; Selection, Projection, Sort and Merge-Join.

Selection

Selection is a very basic operation in database management. It uses relational algebra to select the elements. Now the spec for this operator says that every query is combined with a relational *or*. This means that nothing "fancy" has to be done. It should simply select everytime one or more queries return true.

Projection

The projection is also one of the more basic operations in database management. A projection extracts the columns from a relation, however, unlike in relational algebra, this operator does not eliminate duplicates.

Sort

Sort has to be external. External sorting is used in applications where huge amounts of data has to be sorted, thus the data has to be sorted in chunks since it cannot all be in main memory. A variant of merge-sort will be used, since it can sort on parts of the data and then combine the sorted parts.

Merge-Join

The merge-join assumes that the both inputs are sorted. It then has to merge where possible.

Implementation

Selection.java

Selection starts by assigning local protected variables some values from the parameters.

```
16  public Selection(Iterator iter, Predicate... preds) {  
    this.iterator = iter;  
    this.predicates = preds;  
    this.schema = iter.schema;  
20  this.tuple = null;  
    }
```

The selection process takes place in the hasNext() function. Here it keeps checking if there are more elements to be selected using the evaluate() function. It returns true if it allows the selection and false if there are no more elements to be selected.

```
66  public boolean hasNext() {  
    while(this.iterator.hasNext()) {  
        this.tuple = this.iterator.getNext();  
        for(int i = 0; i < this.predicates.length; ++i) {
```

```

70         if (this.predicates[i].evaluate(this.tuple)) {
            return true;
        }
    }
    return false;
}

```

The getNext() function is what actually gets the elements. It returns a tuple containing the next element. If there are no more elements to be selected, it will throw an exception, telling that there are no more tuples.

```

83 public Tuple getNext() {
    if (this.tuple == null) {
85         throw new IllegalStateException("no_more_tuples");
    } else {
        Tuple tuple = this.tuple;
        this.tuple = null;
        return tuple;
90     }
}

```

Testing

Testing this time around was very successful, assuming it passes in the tests where it returns "failed as expected". It reports that test1, test2 and test3 completed successfully.

Appendix

Pickvictim

```

37 @Override
    public int pickVictim() {
        // Finds the first element in the frametab array, where pin count
        // is equal to zero and returns it.
40         for (int i = 0; i < Minibase.BufferManager.frametab.length;
            i++) {
            if (Minibase.BufferManager.frametab[i].pincnt == 0) {
                return i;
            }
        }
45         // If no pages has zero pins, then it returns -1.
        return -1;
    }

```

bufmgr.java

```

1 package bufmgr;

import java.util.HashMap;

5 import global.GlobalConst;
import global.Minibase;
import global.Page;
import global.PageId;

10 /**

```

```

* <h3>Minibase Buffer Manager</h3> The buffer manager reads disk pages
  ↪ into a
* main memory page as needed. The collection of main memory pages (called
* frames) used by the buffer manager for this purpose is called the buffer
* pool. This is just an array of Page objects. The buffer manager is used
  ↪ by
15 * access methods, heap files, and relational operators to read, write,
* allocate, and de-allocate pages.
*/
@SuppressWarnings("unused")
public class BufMgr implements GlobalConst {
20
    /**
     * Actual pool of pages (can be viewed as an array of byte arrays).
     */
    protected Page[] bufpool;
25
    private boolean debugvalue = false;

    /**
     * Array of descriptors, each containing the pin count, dirty status,
     ↪ etc.
30
     */
    protected FrameDesc[] frametab;

    /**
     * Maps current page numbers to frames; used for efficient lookups.
     */
35    protected HashMap<Integer, FrameDesc> pagemap;

    /**
     * The replacement policy to use.
40
     */
    protected Replacer replacer;

    /**
     * Constructs a buffer manager with the given settings.
45
     * @param numbufs: number of pages in the buffer pool
     */

    public BufMgr(int numbufs) {
50        // initialize the buffer pool and frame table
        bufpool = new Page[numbufs];
        frametab = new FrameDesc[numbufs];
        for (int i = 0; i < numbufs; i++) {
55            bufpool[i] = new Page();
            frametab[i] = new FrameDesc(i);
        }

        // initialize the specialized page map and replacer
        pagemap = new HashMap<Integer, FrameDesc>(numbufs);
60        replacer = new Clock(this);
    }

    /**
     * Allocates a set of new pages, and pins the first one in an
     ↪ appropriate

```

```

65      * frame in the buffer pool.
      *
      * @param firstpg holds the contents of the first page
      * @param run_size number of new pages to allocate
      * @return page id of the first new page
70      * @throws IllegalArgumentException if PIN_MEMCPY and the page is
      *     ↪ pinned
      * @throws IllegalStateException if all pages are pinned (i.e. pool
      *     ↪ exceeded)
      */
      public PageId newPage(Page firstpg, int run_size) {
          // allocate the run
75          PageId firstid = Minibase.DiskManager.allocate_page(run_size);

          // try to pin the first page
          if (debugvalue) {
              System.out.println("trying_to_pin_the_first_page");
80          }
          try {
              pinPage(firstid, firstpg, PIN_MEMCPY);
          } catch (RuntimeException exc) {
              System.out.println("failed_to_pin_the_first_page.");
85              // roll back because pin failed
              for (int i = 0; i < run_size; i++) {
                  firstid.pid += 1;
                  Minibase.DiskManager.deallocate_page(firstid);
              }
90              // re-throw the exception
              throw exc;
          }
          // notify the replacer and return the first new page id
          replacer.newPage(pagemap.get(firstid.pid));
95          return firstid;
      }

      /**
      * Deallocates a single page from disk, freeing it from the pool if
      *     ↪ needed.
100      * Call Minibase.DiskManager.deallocate_page(pageno) to deallocate the
      *     ↪ page before return.
      *
      * @param pageno identifies the page to remove
      * @throws IllegalArgumentException if the page is pinned
      */
105      public void freePage(PageId pageno) throws IllegalArgumentException {
          FrameDesc fdesc = pagemap.get(pageno.pid);
          if (debugvalue){
              System.out.println("freeing_page_with_id_"+pageno.pid);
          }
110          if (fdesc != null) {
              if (fdesc.pincnt != 0) {
                  throw new IllegalArgumentException("The_page_is_pinned.");
              }
              return;
115              //throw new IllegalArgumentException( "page does not excists");
          }
          Minibase.DiskManager.deallocate_page(pageno);
      }
  }

```

```

120  /**
    * Pins a disk page into the buffer pool. If the page is already pinned
    *   ↪ ,
    * this simply increments the pin count. Otherwise, this selects
    *   ↪ another
    * page in the pool to replace, flushing the replaced page to disk if
    * it is dirty.
125  * <p>
    * (If one needs to copy the page from the memory instead of reading
    *   ↪ from
    * the disk, one should set skipRead to PIN_MEMCPY. In this case, the
    *   ↪ page
    * shouldn't be in the buffer pool. Throw an IllegalArgumentException
    *   ↪ if so. )
    *
130  * @param pageno identifies the page to pin
    * @param page if skipread == PIN_MEMCPY, works as an input
    *   ↪ param, holding the contents to be read into the buffer pool
    * if skipread == PIN_DISKIO, works as an output param,
    *   ↪ holding the contents of the pinned page read from the disk
    * @param skipRead PIN_MEMCPY(true) (copy the input page to the buffer
    *   ↪ pool); PIN_DISKIO(false) (read the page from disk)
    * @throws IllegalArgumentException if PIN_MEMCPY and the page is
    *   ↪ pinned
135  * @throws IllegalStateException if all pages are pinned (i.e. pool
    *   ↪ exceeded)
    */
    public void pinPage(PageId pageno, Page page, boolean skipRead) {
        if (debugvalue) System.out.println("pinpage_called_with_pageid_" +
            ↪ pageno.pid + "_Skipread_" + skipRead + "and_page_" + page.
            ↪ toString());

140  // First check if the page is already pinned
        FrameDesc fdesc = Minibase.BufferManager.pagemap.get(pageno
            ↪ .pid);
        if (fdesc != null) {

            // Validate the pin method
145  if (skipRead == PIN_MEMCPY && fdesc.pincnt > 0)
                ↪ throw new IllegalArgumentException(
                    "Page_pinned;_PIN_MEMCPY_not_allowed"
                );
            // Increment pin count, notify the replacer, and wrap the
            ↪ buffer.
            fdesc.pincnt++;
150  replacer.pinPage(fdesc);
            page.setPage(bufpool[fdesc.index]);
            return;
        } // If in pool

155  // Select an available frame
        int frameNo = replacer.pickVictim();
        // If no pages are unpinned, then throw an exception telling that.
        if (frameNo < 0){
            throw new IllegalStateException("All_pages_pinned."
                ↪ );
160  }

```

```

// Pick the frame that is not pinned.
fdesc = Minibase.BufferManager.frameTab[frameNo];
// If the frame was in use and dirty, it should write it to
    ↪ the disk.
165     if( fdesc.pageno.pid != INVALID_PAGEID) {
        flushPage(fdesc.pageno);
        pagemap.remove(fdesc.pageno.pid);
    }

170     //read in the page if requested, and wrap the buffer
    if(skipRead == PIN_MEMCPY) {
        bufpool[frameNo].copyPage(page);
    } else {
        Minibase.DiskManager.read_page(pageno, bufpool[
            ↪ frameNo]);
175     }
    page.setPage(bufpool[frameNo]);

    //update the frame descriptor
    fdesc.pageno.pid = pageno.pid;
180     fdesc.pincnt = 1;
    fdesc.dirty = false;

    // Pin the page and put the updated page in the pagemap.
    pagemap.put(pageno.pid, fdesc);
185     replacer.pinPage(fdesc);
}

/**
 * Unpins a disk page from the buffer pool, decreasing its pin count.
 *
 * @param pageno identifies the page to unpin
 * @param dirty UNPIN_DIRTY if the page was modified, UNPIN_CLEAN
    ↪ otherwise
 * @throws IllegalArgumentException if the page is not present or not
    ↪ pinned
 */
195 public void unpinPage(PageId pageno, boolean dirty) throws
    ↪ IllegalArgumentException {
    if (debugvalue) System.out.println("unpin_page_called_with_pageid"
        ↪ + pageno.pid + "_Dirty_status_" + dirty);
    //Checks if page is dirty.
    // First check if the page is unpinned
    FrameDesc fdesc = pagemap.get(pageno.pid);
200     if (fdesc == null || fdesc.pincnt == 0) throw new
        ↪ IllegalArgumentException(
            "Page_not_pinned;")
    );
    // If dirty, it should write the the page to the disk and then tell
        ↪ that the page is not dirty anymore.
    if(dirty == UNPIN_DIRTY){
205         fdesc.dirty = dirty;
    }
    // Decrement the pin count, since the page is pinned by one less.
        ↪ Also unpin the page and update the page in the
    // pagemap.
    fdesc.pincnt--;

```

```

210     replacer.unpinPage(fdesc);
        //unpin page.

        return;
215     }

    /**
     * Immediately writes a page in the buffer pool to disk, if dirty.
     */
220     public void flushPage(PageId pageno) {
        // Check if page is unpinned
        FrameDesc fdesc = pagemap.get(pageno.pid);
        if (fdesc.dirty == true) {
            // Writes page to disk and sets the dirty-state to false, since
            // ↪ it has not been modified when comparing it
225         // to the same page on the disk.
            Minibase.DiskManager.write_page(fdesc.pageno, bufpool[fdesc.
            ↪ index]);
            fdesc.dirty = false;
            pagemap.put(pageno.pid, fdesc);
        }
230     }

    /**
     * Immediately writes all dirty pages in the buffer pool to disk.
     */
235     public void flushAllPages() {
        for (int i = 0 ; i < Minibase.BufferManager.frametab.length; i
            ↪ ++ ){
            if (debugvalue) {
                System.out.println("flushing_page_" + Minibase.
                ↪ BufferManager.frametab[i].pageno.pid);
            }
240         if (Minibase.BufferManager.frametab[i].pageno.pid > 0) {
            flushPage(Minibase.BufferManager.frametab[i].pageno);
        }
    }
245 }

    /**
     * Gets the total number of buffer frames.
     */
    public int getNumBuffers() {
250         return Minibase.BufferManager.bufpool.length;
    }

    /**
     * Gets the total number of unpinned buffer frames.
255     */
    public int getNumUnpinned() {
        // Using a loop, this checks the state of each frame. Each time
        // ↪ an unpinned frame is found, "j" is incremented.
        // In the end "j" is returned, as that must be the total amount of
        // ↪ unpinned buffer frames.
        int j = 0;
260         for (int i = 0 ; i < Minibase.BufferManager.frametab.length; i++ )
            ↪ {

```


265

```
        if (0 == Minibase.BufferManager.frametab[i].pincnt) j++;
    }
    return j;
}

} // public class BufMgr implements GlobalConst
```